

Subject: A new kind of architecture review — and an ask for your team

Team,

Most of you know how Q4 went. Two P2 incidents, one audit finding, and a stack of late-night bridge calls I'd rather not repeat. The common thread in the post-mortems and audit reports was clear: the signs were in the architecture docs, the manifests, and the policy library; just never in the same place at the same time.

Meanwhile, the group board keeps asking what we're actually doing with AI internally. Fair question.

So here's an experiment, and an ask.

We have a service called **QuoteCraft** due for its annual architecture review. Silver criticality. Policy quotation microservice, running on AlphaPaaS, with a mix of AWS and Azure managed services behind it. I want each of your teams to build an **agentic AI system** that can perform that review — across cost, security, scalability, performance and availability, grounded in our own policies, with real recommendations, not generic ones.

This isn't a demo. If what you build is any good, the Cloud Broker teams will use it on the post Atlas journey.

A few ground rules:

- You'll be assigned AWS or Azure to build on. The point is to come out of this with real depth in the agentic tooling each vendor offers, Bedrock AgentCore or AI Foundry, because those are our two main cloud service providers, and we need people who've touched these services in anger.
- Use Python and LaghChain / LangGraph to build your agents, we want to be able to have a common ecosystem so that we can derive best practices and reusable assets from each-other.
- Be frugal! We need to build expertise not break the bank. Be particularly mindful to how many LLM tokens you consume.
- **No scoring, no winners.** At the end of the month we present to each other, compare approaches, and keep what works. I'd rather end with five different ideas than one winner.
- Use the “**Cloud Agentic AI Hackathon – May**” Teams channel for cross-team questions. AWS and Azure will each run a deep-dive session in week one; office hours follow in weeks two and three; we close with a two-day sprint on 27th - 28th May (remote and in Enjoy, for those who are Paris-based).

Kickoff Monday. Materials pack: intake form, repo link, policy library, pricing API spec, all on the hackathon landing page. **Read the intake form first.** It's the contract.

Go build something good.

— Sophie

Sophie Laurent Chief Technology Officer, AlphaInsure